

Machine Learning in Precision Medicine

Linear Regression & Regularization

Prof. Dr. Christoph Lippert

Dr. Stefan Konigorski

Digital Health & Machine Learning

Formalities

All lecture and exercise materials are now available at
<https://open.hpi.de/courses/MLPrecisionMedicineSoSe2019>.

Further news and information through our mailing list, where you also submit your exercise solutions.

Further questions regarding formalities?

Contents of today: Linear Regression

- Introduction
- Maximum likelihood estimation
- Polynomial curve fitting
- Solutions to overfitting
- Regularizes least squares

Linear Regression Recap & Quiz





What is a linear (regression) model?

Linear Regression Recap & Quiz





What is a linear (regression) model?



Linear Regression Recap & Quiz



What is a linear (regression) model?



→ linear in parameters (e.g. x^3 could be redefined as z).

Linear Regression Recap & Quiz



What is the likelihood?

Linear Regression Recap & Quiz



What is the likelihood?

Probability of data given parameters: $p(\mathcal{D}|\boldsymbol{\theta})$

What is the likelihood?

Probability of data given parameters: $p(\mathcal{D}|\theta)$

What is the rationale behind using maximum likelihood as a criterion for finding the best line fitting 2D data points?

What is the likelihood?

Probability of data given parameters: $p(\mathcal{D}|\theta)$

What is the rationale behind using maximum likelihood as a criterion for finding the best line fitting 2D data points?

How do you obtain the maximum likelihood values for your parameters (= MLE, maximum likelihood estimates)?

What is the likelihood?

Probability of data given parameters: $p(\mathcal{D}|\theta)$

What is the rationale behind using maximum likelihood as a criterion for finding the best line fitting 2D data points?

How do you obtain the maximum likelihood values for your parameters (= MLE, maximum likelihood estimates)?

Set first derivative (w.r.t. each parameter of interest) of the log-likelihood to 0, solve.

What do you need to know to write out the likelihood?

What do you need to know to write out the likelihood?

- Model
- Distribution (pdf) of model, whether points are iid (i.e. whether pdf of all points factorizes into product of single points)
- and ...

What do you need to know to write out the likelihood?

- Model
- Distribution (pdf) of model, whether points are iid (i.e. whether pdf of all points factorizes into product of single points)
- and ...
- that's all! Easy to get closed form, once you know (i.e. assume) a specific parametric distribution.

Linear Regression

Maximum likelihood



- Setting: data \mathcal{D} consisting of N data points $\{\mathbf{x}_n, y_n\}_{n=1}^N$, where \mathbf{x}_n contains D features $\mathbf{x}_n = (x_{n1}, \dots, x_{nD})$.

- Setting: data \mathcal{D} consisting of N data points $\{\mathbf{x}_n, y_n\}_{n=1}^N$, where \mathbf{x}_n contains D features $\mathbf{x}_n = (x_{n1}, \dots, x_{nD})$.

- In matrix formulation: $\mathbf{X} = \begin{bmatrix} x_{11} & \dots & x_{1D} \\ \vdots & \ddots & \vdots \\ x_{N1} & \dots & x_{ND} \end{bmatrix} = \begin{bmatrix} \mathbf{x}_1 \\ \vdots \\ \mathbf{x}_N \end{bmatrix}.$

- Setting: data \mathcal{D} consisting of N data points $\{\mathbf{x}_n, y_n\}_{n=1}^N$, where \mathbf{x}_n contains D features $\mathbf{x}_n = (x_{n1}, \dots, x_{nD})$.

- In matrix formulation: $\mathbf{X} = \begin{bmatrix} x_{11} & \dots & x_{1D} \\ \vdots & \ddots & \vdots \\ x_{N1} & \dots & x_{ND} \end{bmatrix} = \begin{bmatrix} \mathbf{x}_1 \\ \vdots \\ \mathbf{x}_N \end{bmatrix}$.

- Assumption: linear model $y_n = \mathbf{x}_n \cdot \boldsymbol{\beta} + \epsilon_n$ with $\boldsymbol{\beta} \in \mathbb{R}^D$, $\epsilon_n \sim \mathcal{N}(0, \sigma^2)$.

- Setting: data \mathcal{D} consisting of N data points $\{\mathbf{x}_n, y_n\}_{n=1}^N$, where \mathbf{x}_n contains D features $\mathbf{x}_n = (x_{n1}, \dots, x_{nD})$.

- In matrix formulation:
$$\mathbf{X} = \begin{bmatrix} x_{11} & \dots & x_{1D} \\ \vdots & \ddots & \vdots \\ x_{N1} & \dots & x_{ND} \end{bmatrix} = \begin{bmatrix} \mathbf{x}_1 \\ \vdots \\ \mathbf{x}_N \end{bmatrix}.$$

- Assumption: linear model $y_n = \mathbf{x}_n \cdot \boldsymbol{\beta} + \epsilon_n$ with $\boldsymbol{\beta} \in \mathbb{R}^D$, $\epsilon_n \sim \mathcal{N}(0, \sigma^2)$.
- Aim: Learn parameters by maximizing the likelihood.

- Setting: data \mathcal{D} consisting of N data points $\{\mathbf{x}_n, y_n\}_{n=1}^N$, where \mathbf{x}_n contains D features $\mathbf{x}_n = (x_{n1}, \dots, x_{nD})$.

- In matrix formulation:
$$\mathbf{X} = \begin{bmatrix} x_{11} & \dots & x_{1D} \\ \vdots & \ddots & \vdots \\ x_{N1} & \dots & x_{ND} \end{bmatrix} = \begin{bmatrix} \mathbf{x}_1 \\ \vdots \\ \mathbf{x}_N \end{bmatrix}.$$

- Assumption: linear model $y_n = \mathbf{x}_n \cdot \boldsymbol{\beta} + \epsilon_n$ with $\boldsymbol{\beta} \in \mathbb{R}^D$, $\epsilon_n \sim \mathcal{N}(0, \sigma^2)$.
- Aim: Learn parameters by maximizing the likelihood.
- Just to make sure: what are the parameters to be learned?

- Taking the logarithm, we obtain

$$\begin{aligned}\ln p(\mathbf{y}|\boldsymbol{\theta}, \sigma^2) &= \sum_{n=1}^N \ln \mathcal{N}(y_n \mid \mathbf{x}_n \cdot \boldsymbol{\beta}, \sigma^2) \\ &= -\frac{N}{2} \ln 2\pi\sigma^2 - \frac{1}{2\sigma^2} \underbrace{\sum_{n=1}^N (y_n - \mathbf{x}_n \cdot \boldsymbol{\beta})^2}_{\text{Sum of squares}}\end{aligned}$$

- Taking the logarithm, we obtain

$$\begin{aligned}\ln p(\mathbf{y}|\boldsymbol{\theta}, \sigma^2) &= \sum_{n=1}^N \ln \mathcal{N}(y_n \mid \mathbf{x}_n \cdot \boldsymbol{\beta}, \sigma^2) \\ &= -\frac{N}{2} \ln 2\pi\sigma^2 - \frac{1}{2\sigma^2} \underbrace{\sum_{n=1}^N (y_n - \mathbf{x}_n \cdot \boldsymbol{\beta})^2}_{\text{Sum of squares}}\end{aligned}$$

- The likelihood is **maximized** when the **squared error** is **minimized**.

- Taking the logarithm, we obtain

$$\begin{aligned}\ln p(\mathbf{y}|\boldsymbol{\theta}, \sigma^2) &= \sum_{n=1}^N \ln \mathcal{N}(y_n \mid \mathbf{x}_n \cdot \boldsymbol{\beta}, \sigma^2) \\ &= -\frac{N}{2} \ln 2\pi\sigma^2 - \frac{1}{2\sigma^2} \underbrace{\sum_{n=1}^N (y_n - \mathbf{x}_n \cdot \boldsymbol{\beta})^2}_{\text{Sum of squares}}\end{aligned}$$

- The likelihood is **maximized** when the **squared error** is **minimized**.
- **Least squares** and maximum likelihood are equivalent.

- Taking the logarithm, we obtain

$$\begin{aligned}\ln p(\mathbf{y}|\boldsymbol{\theta}, \sigma^2) &= \sum_{n=1}^N \ln \mathcal{N}(y_n \mid \mathbf{x}_n \cdot \boldsymbol{\beta}, \sigma^2) \\ &= -\frac{N}{2} \ln 2\pi\sigma^2 - \frac{1}{2\sigma^2} \underbrace{\sum_{n=1}^N (y_n - \mathbf{x}_n \cdot \boldsymbol{\beta})^2}_{\text{Sum of squares}}\end{aligned}$$

- The likelihood is **maximized** when the **squared error** is **minimized**.
- **Least squares** and maximum likelihood are equivalent.
- Question: which assumption do you make using least squares to learn parameter values?


- Taking the logarithm, we obtain

$$\begin{aligned}\ln p(\mathbf{y}|\boldsymbol{\theta}, \sigma^2) &= \sum_{n=1}^N \ln \mathcal{N}(y_n | \mathbf{x}_n \cdot \boldsymbol{\beta}, \sigma^2) \\ &= -\frac{N}{2} \ln 2\pi\sigma^2 - \frac{1}{2\sigma^2} \underbrace{\sum_{n=1}^N (y_n - \mathbf{x}_n \cdot \boldsymbol{\beta})^2}_{\text{Sum of squares}}\end{aligned}$$

- The likelihood is **maximized** when the **squared error** is **minimized**.
- **Least squares** and maximum likelihood are equivalent.
- Question: which assumption do you make using least squares to learn parameter values?
- i.e. MLEs under linear model are (relatively) robust against wrong distribution assumption!

Linear Regression


Maximum Likelihood and Least Squares



$$\operatorname{argmin}_{\beta} \frac{1}{2} \sum_{n=1}^N (y_n - \mathbf{x}_n \cdot \beta)^2$$

Linear Regression

Maximum Likelihood and Least Squares



How would the estimated regression line change?

Linear Regression

Derive parameter values



- **Derivative** w.r.t. a single weight entry β_i

$$\frac{d}{d\beta_d} \ln p(\mathbf{y}|\boldsymbol{\beta}, \sigma^2) = \frac{d}{d\beta_d} \left[-\frac{N}{2} \ln 2\pi\sigma^2 - \frac{1}{2\sigma^2} \sum_{n=1}^N (y_n - \mathbf{x}_n \cdot \boldsymbol{\beta})^2 \right]$$

Linear Regression

Derive parameter values



- **Derivative** w.r.t. a single weight entry β_i

$$\begin{aligned}\frac{d}{d\beta_d} \ln p(\mathbf{y}|\boldsymbol{\beta}, \sigma^2) &= \frac{d}{d\beta_d} \left[-\frac{N}{2} \ln 2\pi\sigma^2 - \frac{1}{2\sigma^2} \sum_{n=1}^N (y_n - \mathbf{x}_n \cdot \boldsymbol{\beta})^2 \right] \\ &= \left[\frac{1}{\sigma^2} \sum_{n=1}^N x_{nd}(y_n - \mathbf{x}_n \cdot \boldsymbol{\beta}) \right]\end{aligned}$$

Linear Regression

Derive parameter values



- **Derivative** w.r.t. a single weight entry β_i

$$\begin{aligned}\frac{d}{d\beta_d} \ln p(\mathbf{y}|\boldsymbol{\beta}, \sigma^2) &= \frac{d}{d\beta_d} \left[-\frac{N}{2} \ln 2\pi\sigma^2 - \frac{1}{2\sigma^2} \sum_{n=1}^N (y_n - \mathbf{x}_n \cdot \boldsymbol{\beta})^2 \right] \\ &= \left[\frac{1}{\sigma^2} \sum_{n=1}^N x_{nd}(y_n - \mathbf{x}_n \cdot \boldsymbol{\beta}) \right]\end{aligned}$$

- Set **gradient** w.r.t. $\boldsymbol{\beta}$ to zero [vector holding the derivatives $\forall \beta_d$]

$$\nabla_{\boldsymbol{\beta}} \ln p(\mathbf{y}|\boldsymbol{\beta}, \sigma^2) =$$

Linear Regression

Derive parameter values



- **Derivative** w.r.t. a single weight entry β_i

$$\begin{aligned}\frac{d}{d\beta_d} \ln p(\mathbf{y}|\boldsymbol{\beta}, \sigma^2) &= \frac{d}{d\beta_d} \left[-\frac{N}{2} \ln 2\pi\sigma^2 - \frac{1}{2\sigma^2} \sum_{n=1}^N (y_n - \mathbf{x}_n \cdot \boldsymbol{\beta})^2 \right] \\ &= \left[\frac{1}{\sigma^2} \sum_{n=1}^N x_{nd}(y_n - \mathbf{x}_n \cdot \boldsymbol{\beta}) \right]\end{aligned}$$

- Set **gradient** w.r.t. $\boldsymbol{\beta}$ to zero [vector holding the derivatives $\forall \beta_d$]

$$\nabla_{\boldsymbol{\beta}} \ln p(\mathbf{y}|\boldsymbol{\beta}, \sigma^2) = \frac{1}{\sigma^2} \sum_{n=1}^N \mathbf{x}_n^T (y_n - \mathbf{x}_n \cdot \boldsymbol{\beta}) = \mathbf{0} \quad (\text{where } \mathbf{0} \text{ is a vector of 0s})$$

Linear Regression

Derive parameter values



- Set **gradient** w.r.t. β to zero [vector holding the derivatives $\forall \beta_d$]

$$\nabla_{\beta} \ln p(\mathbf{y}|\beta, \sigma^2) = \frac{1}{\sigma^2} \sum_{n=1}^N \mathbf{x}_n^T (y_n - \mathbf{x}_n \cdot \beta) = \mathbf{0} \quad (\text{where } \mathbf{0} \text{ is a vector of 0s})$$

Linear Regression

Derive parameter values



- Set **gradient** w.r.t. β to zero [vector holding the derivatives $\forall \beta_d$]

$$\nabla_{\beta} \ln p(\mathbf{y}|\beta, \sigma^2) = \frac{1}{\sigma^2} \sum_{n=1}^N \mathbf{x}_n^T (y_n - \mathbf{x}_n \cdot \beta) = \mathbf{0} \quad (\text{where } \mathbf{0} \text{ is a vector of 0s})$$

- In matrix formulation: $\frac{1}{\sigma^2} \mathbf{X}^T (\mathbf{y} - \mathbf{X} \cdot \beta) = \mathbf{0}$.

Linear Regression

Derive parameter values



- Set **gradient** w.r.t. β to zero [vector holding the derivatives $\forall \beta_d$]

$$\nabla_{\beta} \ln p(\mathbf{y}|\beta, \sigma^2) = \frac{1}{\sigma^2} \sum_{n=1}^N \mathbf{x}_n^T (y_n - \mathbf{x}_n \cdot \beta) = \mathbf{0} \quad (\text{where } \mathbf{0} \text{ is a vector of 0s})$$

- In matrix formulation: $\frac{1}{\sigma^2} \mathbf{X}^T (\mathbf{y} - \mathbf{X} \cdot \beta) = \mathbf{0}$.
- Solving for β yields: $\beta_{\text{ML}} =$

Linear Regression

Derive parameter values



- Set **gradient** w.r.t. β to zero [vector holding the derivatives $\forall \beta_d$]

$$\nabla_{\beta} \ln p(\mathbf{y}|\beta, \sigma^2) = \frac{1}{\sigma^2} \sum_{n=1}^N \mathbf{x}_n^T (y_n - \mathbf{x}_n \cdot \beta) = \mathbf{0} \quad (\text{where } \mathbf{0} \text{ is a vector of 0s})$$

- In matrix formulation: $\frac{1}{\sigma^2} \mathbf{X}^T (\mathbf{y} - \mathbf{X} \cdot \beta) = \mathbf{0}$.
- Solving for β yields: $\beta_{\text{ML}} = \underbrace{(\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T}_{\text{Pseudo inverse of } \mathbf{X}} \mathbf{y}$

Linear Regression

Derive parameter values



- Set **gradient** w.r.t. β to zero [vector holding the derivatives $\forall \beta_d$]

$$\nabla_{\beta} \ln p(\mathbf{y}|\beta, \sigma^2) = \frac{1}{\sigma^2} \sum_{n=1}^N \mathbf{x}_n^T (y_n - \mathbf{x}_n \cdot \beta) = \mathbf{0} \quad (\text{where } \mathbf{0} \text{ is a vector of 0s})$$

- In matrix formulation: $\frac{1}{\sigma^2} \mathbf{X}^T (\mathbf{y} - \mathbf{X} \cdot \beta) = \mathbf{0}$.
- Solving for β yields: $\beta_{\text{ML}} = \underbrace{(\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T}_{\text{Pseudo inverse of } \mathbf{X}} \mathbf{y}$
- Question: what happened to $\frac{1}{\sigma^2}$?

Linear Regression

Derive parameter values



- Set **gradient** w.r.t. β to zero [vector holding the derivatives $\forall \beta_d$]

$$\nabla_{\beta} \ln p(\mathbf{y}|\beta, \sigma^2) = \frac{1}{\sigma^2} \sum_{n=1}^N \mathbf{x}_n^T (y_n - \mathbf{x}_n \cdot \beta) = \mathbf{0} \quad (\text{where } \mathbf{0} \text{ is a vector of 0s})$$

- In matrix formulation: $\frac{1}{\sigma^2} \mathbf{X}^T (\mathbf{y} - \mathbf{X} \cdot \beta) = \mathbf{0}$.
- Solving for β yields: $\beta_{\text{ML}} = \underbrace{(\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T}_{\text{Pseudo inverse of } \mathbf{X}} \mathbf{y}$
- Question: what happened to $\frac{1}{\sigma^2}$?
- Learn parameter values for σ^2 similarly.

Linear Regression

Polynomial curve fitting



- The above can be extended to more general functions of \mathbf{x}_n that are linear in β . For simplicity, focus here on $\mathbf{x}_n = x$ (i.e. one feature).

For example:

Linear Regression

Polynomial curve fitting



- The above can be extended to more general functions of \mathbf{x}_n that are linear in β . For simplicity, focus here on $\mathbf{x}_n = x$ (i.e. one feature).

For example:

- Use the polynomials up to degree M to construct new features

$$\begin{aligned}f(x, \beta) &= \beta_0 + \beta_1 x + \beta_2 x^2 + \dots + \beta_M x^M \\ &= \phi(x) \cdot \beta\end{aligned}$$

with $\phi(x) = (1, x, x^2, \dots, x^M)$, $\phi_m(x) = x^m$, $\beta = (\beta_0, \beta_1, \dots, \beta_M)^T \in \mathbb{R}^M$.

Linear Regression

Polynomial curve fitting



- The above can be extended to more general functions of \mathbf{x}_n that are linear in β . For simplicity, focus here on $\mathbf{x}_n = x$ (i.e. one feature).

For example:

- Use the polynomials up to degree M to construct new features

$$\begin{aligned}f(x, \beta) &= \beta_0 + \beta_1 x + \beta_2 x^2 + \dots + \beta_M x^M \\ &= \phi(x) \cdot \beta\end{aligned}$$

with $\phi(x) = (1, x, x^2, \dots, x^M)$, $\phi_m(x) = x^m$, $\beta = (\beta_0, \beta_1, \dots, \beta_M)^T \in \mathbb{R}^M$.

- Then, these polynomials can be used to predict \mathbf{y} with $\mathbf{y} = \phi(x) \cdot \beta + \epsilon$.

Linear Regression

Polynomial curve fitting



- The above can be extended to more general functions of \mathbf{x}_n that are linear in β . For simplicity, focus here on $\mathbf{x}_n = x$ (i.e. one feature).

For example:

- Use the polynomials up to degree M to construct new features

$$\begin{aligned}f(x, \beta) &= \beta_0 + \beta_1 x + \beta_2 x^2 + \dots + \beta_M x^M \\ &= \phi(x) \cdot \beta\end{aligned}$$


with $\phi(x) = (1, x, x^2, \dots, x^M)$, $\phi_m(x) = x^m$, $\beta = (\beta_0, \beta_1, \dots, \beta_M)^T \in \mathbb{R}^M$.

- Then, these polynomials can be used to predict \mathbf{y} with $\mathbf{y} = \phi(x) \cdot \beta + \epsilon$.
- Similarly, ϕ can be any feature mapping (\rightarrow call ϕ_m basis functions, see next lecture).

Linear Regression

Polynomial curve fitting

The degree M of the polynomial is crucial.




(C.M. Bishop, Pattern Recognition and Machine Learning)

Linear Regression

Polynomial curve fitting

The degree M of the polynomial is crucial.




(C.M. Bishop, Pattern Recognition and Machine Learning)

Linear Regression

Polynomial curve fitting

The degree M of the polynomial is crucial.




(C.M. Bishop, Pattern Recognition and Machine Learning)

Linear Regression

Polynomial curve fitting

The degree M of the polynomial is crucial.



(C.M. Bishop, Pattern Recognition and Machine Learning)

Question: So how should you choose the best M ?

Question: So how should you choose the best M ?

Idea: Fit regression model for any M , compute mean squared error (MSE), choose M with lowest MSE.

Question: So how should you choose the best M ?

Idea: Fit regression model for any M , compute mean squared error (MSE), choose M with lowest MSE.

Problems?

Question: So how should you choose the best M ?

Idea: Fit regression model for any M , compute mean squared error (MSE), choose M with lowest MSE.

Problems? Overfitting!

Linear Regression


Train vs. Test error



- Split sample in training and test set.
- Choose M based on test error.

Linear Regression Train vs. Test error

- Split sample in training and test set.
- Choose M based on test error.




(C.M. Bishop, Pattern Recognition and Machine Learning)

- Intuition? Interpretation?

Linear Regression Train vs. Test error

How do you compute the test error?

- Split sample in training and test set.
- Choose M based on test error.



(C.M. Bishop, Pattern Recognition and Machine Learning)


- Intuition? Interpretation?

Linear Regression Train vs. Test error

- Split sample in training and test set.
- Choose M based on test error.

How do you compute the test error?

Cross validation:




(C.M. Bishop, Pattern Recognition and Machine Learning)

- Intuition? Interpretation?

Linear Regression

Train vs. Test error

- Split sample in training and test set.
- Choose M based on test error.




(C.M. Bishop, Pattern Recognition and Machine Learning)

- Intuition? Interpretation?

How do you compute the test error?

Cross validation:



- randomly assign samples (\mathbf{x}_n, y_n) to S sets of equal size
- for each set $s \in S$ (e.g. $S = 4$) and $m \in [1, \dots, M]$:
 - train model on $S - 1$ remaining sets
 - predict on s and compute MSE or root MSE E_{RMS} .
- compute average MSE for degree m .
- pick m with lowest MSE.





(C.M. Bishop, Pattern Recognition and Machine Learning)



What else can you do to improve model, i.e. learn better weights, find better curve fit, and be less prone to overfitting?

Linear Regression Get more data





Others ways to avoid overfitting and fit complex model for limited number of observations?



Others ways to avoid overfitting and fit complex model for limited number of observations?

Regularization of weights

Regularize the regression weights β .

Loss function:

$$E_2(\beta) = \underbrace{\frac{1}{2} \sum_{n=1}^N (y_n - \mathbf{x}_n \cdot \beta)^2}_{\text{squared error}} + \underbrace{\frac{\lambda}{2} \sum_{d=1}^D \beta_d^2}_{l_2\text{-norm regularizer}}$$

where $\|\beta\|^2 = \sqrt{\sum_{d=1}^D \beta_d^2}$ is the l_2 -norm of β .

Regularize the regression weights β .

Loss function:
$$E_2(\beta) = \underbrace{\frac{1}{2} \sum_{n=1}^N (y_n - \mathbf{x}_n \cdot \beta)^2}_{\text{squared error}} + \underbrace{\frac{\lambda}{2} \sum_{d=1}^D \beta_d^2}_{l_2\text{-norm regularizer}}$$

where $\|\beta\|^2 = \sqrt{\sum_{d=1}^D \beta_d^2}$ is the l_2 -norm of β .

What effect does this have? How will the learned weights be different?

Regularize the regression weights β .

Loss function:

$$E_2(\beta) = \underbrace{\frac{1}{2} \sum_{n=1}^N (y_n - \mathbf{x}_n \cdot \beta)^2}_{\text{squared error}} + \underbrace{\frac{\lambda}{2} \sum_{d=1}^D \beta_d^2}_{l_2\text{-norm regularizer}}$$

where $\|\beta\|^2 = \sqrt{\sum_{d=1}^D \beta_d^2}$ is the l_2 -norm of β .



What effect does this have? How will the learned weights be different?

- Penalizes large weights.
- Reduces the complexity of the function that associates \mathbf{x} with \mathbf{y} , i.e. learn parsimonious model.
- Also known as shrinkage methods.

Linear Regression Regularization

Loss function:

$$E_2(\beta) = \underbrace{\frac{1}{2} \sum_{n=1}^N (y_n - \mathbf{x}_n \cdot \beta)^2}_{\text{squared error}} + \underbrace{\frac{\lambda}{2} \sum_{d=1}^D \beta_d^2}_{l_2\text{-norm regularizer}}$$



$$E_2(\beta) = \underbrace{\frac{1}{2} \sum_{n=1}^N (y_n - \mathbf{x}_n \cdot \beta)^2}_{\text{squared error}} + \underbrace{\frac{\lambda}{2} \sum_{d=1}^D \beta_d^2}_{l_2\text{-norm regularizer}}$$

Quiz: How to choose an optimal λ ?

$$E_2(\beta) = \underbrace{\frac{1}{2} \sum_{n=1}^N (y_n - \mathbf{x}_n \cdot \beta)^2}_{\text{squared error}} + \underbrace{\frac{\lambda}{2} \sum_{d=1}^D \beta_d^2}_{l_2\text{-norm regularizer}}$$


Quiz: How to choose an optimal λ ?

Answer: Look at the test error!

$$E_2(\beta) = \underbrace{\frac{1}{2} \sum_{n=1}^N (y_n - \mathbf{x}_n \cdot \beta)^2}_{\text{squared error}} + \underbrace{\frac{\lambda}{2} \sum_{d=1}^D \beta_d^2}_{l_2\text{-norm regularizer}}$$

Quiz: How to choose an optimal λ ?

Answer: Look at the test error!



Linear Regression

Solving ridge regression



Minimize the loss function $E_2(\beta) = \underbrace{\frac{1}{2} \sum_{n=1}^N (y_n - \mathbf{x}_n \cdot \beta)^2}_{\text{squared error}} + \underbrace{\frac{\lambda}{2} \sum_{d=1}^D \beta_d^2}_{l_2\text{-norm regularizer}}$

Linear Regression

Solving ridge regression

Minimize the loss function $E_2(\beta) = \underbrace{\frac{1}{2} \sum_{n=1}^N (y_n - \mathbf{x}_n \cdot \beta)^2}_{\text{squared error}} + \underbrace{\frac{\lambda}{2} \sum_{d=1}^D \beta_d^2}_{l_2\text{-norm regularizer}}$

- **Derivative** w.r.t. a single weight entry β_i

$$\frac{d}{d\beta_d} \left[\frac{1}{2} \sum_{n=1}^N (y_n - \mathbf{x}_n \cdot \beta)^2 + \frac{\lambda}{2} \sum_{d=1}^D \beta_d^2 \right]$$

Linear Regression

Solving ridge regression



Minimize the loss function $E_2(\beta) = \underbrace{\frac{1}{2} \sum_{n=1}^N (y_n - \mathbf{x}_n \cdot \beta)^2}_{\text{squared error}} + \underbrace{\frac{\lambda}{2} \sum_{d=1}^D \beta_d^2}_{l_2\text{-norm regularizer}}$

- **Derivative** w.r.t. a single weight entry β_i

$$\frac{d}{d\beta_d} \left[\frac{1}{2} \sum_{n=1}^N (y_n - \mathbf{x}_n \cdot \beta)^2 + \frac{\lambda}{2} \sum_{d=1}^D \beta_d^2 \right] = \sum_{n=1}^N -x_{nd}(y_n - \mathbf{x}_n \cdot \beta) + \lambda\beta_d$$

Linear Regression

Solving ridge regression



Minimize the loss function $E_2(\beta) = \underbrace{\frac{1}{2} \sum_{n=1}^N (y_n - \mathbf{x}_n \cdot \beta)^2}_{\text{squared error}} + \underbrace{\frac{\lambda}{2} \sum_{d=1}^D \beta_d^2}_{l_2\text{-norm regularizer}}$

- **Derivative** w.r.t. a single weight entry β_i

$$\frac{d}{d\beta_d} \left[\frac{1}{2} \sum_{n=1}^N (y_n - \mathbf{x}_n \cdot \beta)^2 + \frac{\lambda}{2} \sum_{d=1}^D \beta_d^2 \right] = \sum_{n=1}^N -x_{nd}(y_n - \mathbf{x}_n \cdot \beta) + \lambda\beta_d$$

- Set **gradient** w.r.t. β to zero [vector holding the derivatives $\forall \beta_d$]

$$\nabla_{\beta} E_2(\beta) = \sum_{n=1}^N \mathbf{X}^T (\mathbf{x}_n \cdot \beta - y_n) + \lambda \beta = \mathbf{0}$$

Linear Regression

Solving ridge regression

Minimize the loss function $E_2(\beta) = \underbrace{\frac{1}{2} \sum_{n=1}^N (y_n - \mathbf{x}_n \cdot \beta)^2}_{\text{squared error}} + \underbrace{\frac{\lambda}{2} \sum_{d=1}^D \beta_d^2}_{l_2\text{-norm regularizer}}$

- **Derivative** w.r.t. a single weight entry β_i


$$\frac{d}{d\beta_d} \left[\frac{1}{2} \sum_{n=1}^N (y_n - \mathbf{x}_n \cdot \beta)^2 + \frac{\lambda}{2} \sum_{d=1}^D \beta_d^2 \right] = \sum_{n=1}^N -x_{nd}(y_n - \mathbf{x}_n \cdot \beta) + \lambda\beta_d$$

- Set **gradient** w.r.t. β to zero [vector holding the derivatives $\forall \beta_d$]

$$\begin{aligned} \nabla_{\beta} E_2(\beta) &= \sum_{n=1}^N \mathbf{X}^T (\mathbf{x}_n \cdot \beta - y_n) + \lambda \beta = \mathbf{0} \\ \implies \beta_{Ridge} &= (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^T \mathbf{y} \end{aligned}$$

A more general regularization:

$$E_q(\beta) = \underbrace{\frac{1}{2} \sum_{n=1}^N (y_n - \mathbf{x}_n \cdot \beta)^2}_{\text{squared error}} + \underbrace{\frac{\lambda}{2} \sum_{d=1}^D |\beta_d|^q}_{l_q\text{-norm regularizer}}$$



(C.M. Bishop, Pattern Recognition and Machine Learning)

A more general regularization:

$$E_q(\beta) = \underbrace{\frac{1}{2} \sum_{n=1}^N (y_n - \mathbf{x}_n \cdot \beta)^2}_{\text{squared error}} + \underbrace{\frac{\lambda}{2} \sum_{d=1}^D |\beta_d|^q}_{l_q\text{-norm regularizer}}$$

- l_1 : Lasso \rightarrow feature selection, sparse solution
- l_2 : Ridge \rightarrow shrinks weights to 0
- Linear combination of l_1 and l_2 : Elastic net

- Maximum likelihood in linear regression
- Polynomial linear regression
- Cross-classification
- Regularized regression with l_2 -norm (ridge regression)

Questions?